

Literature Cited

References for the poster *Designing and Evaluating AI-Assisted Deductive Research Analysis Workflow for Education Impact Studies*.

Graham Chickering, Billie Day and Garry Davis

- Abramson, C. M., Prendergast, T., Li, Z., & Dohan, D. (2026). Qualitative research in an era of artificial intelligence: A pragmatic approach to data analysis, workflow, and computation. *Annual Review of Sociology*, 52(1), 35–61. <https://doi.org/10.1146/annurev-soc-011824-104836>
- Ashwin, J., Chhabra, A., & Rao, V. (2026). Using large language models for qualitative analysis can introduce serious bias. *Sociological Methods & Research*, 55(3), 795–839. <https://doi.org/10.1177/00491241251338246>
- Barrie, C., Palaiologou, E., & Törnberg, P. (2024). *Prompt stability scoring for text annotation with large language models* (arXiv:2407.02039). arXiv. <https://arxiv.org/abs/2407.02039>
- Chew, R., Bollenbacher, J., Wenger, M., Speer, J., & Kim, A. (2023). *LLM-assisted content analysis: Using large language models to support deductive coding* (arXiv:2306.14924). arXiv. <https://arxiv.org/abs/2306.14924>
- DiGiuseppe, M. R., & Flynn, M. E. (2026). Scaling open-ended survey responses using LLM-paired comparisons. *Public Opinion Quarterly*, 90(3), 630–656. <https://doi.org/10.1093/poq/nfag013>
- Dunivin, Z. O. (2025). Scaling hermeneutics: A guide to qualitative coding with LLMs for reflexive content analysis. *EPJ Data Science*, 14(1), Article 28. <https://doi.org/10.1140/epjds/s13688-025-00548-8>
- Feuston, J. L., & Brubaker, J. R. (2021). Putting tools in their place: The role of time and perspective in human-AI collaboration for qualitative analysis. *Proceedings of the ACM on Human-Computer Interaction*, 5(CSCW2), 1–25. <https://doi.org/10.1145/3479856>
- Godfrey, C., Nie, P., Ostapuk, N., Ken, D., Gao, S., & Inati, S. (2025). *Likert or not: LLM absolute relevance judgments on fine-grained ordinal scales* (arXiv:2505.19334). arXiv. <https://arxiv.org/abs/2505.19334>
- Haldar, R., & Hockenmaier, J. (2025). Rating roulette: Self-inconsistency in LLM-as-a-judge frameworks. In *Findings of the Association for Computational Linguistics: EMNLP 2025* (pp. 24986–25004). Association for Computational Linguistics. <https://aclanthology.org/2025.findings-emnlp.1361/>
- Hayes, A. F., & Krippendorff, K. (2007). Answering the call for a standard reliability measure for coding data. *Communication Methods and Measures*, 1(1), 77–89. <https://doi.org/10.1080/19312450709336664>
- Jones, K. M. L. (2025). Generative AI in qualitative research and related transparency problems: A novel heuristic for disclosing uses of AI. *International Journal of Qualitative Methods*, 24, Article 16094069251404329. <https://doi.org/10.1177/16094069251404329>

- Kim, S., Shin, J., Cho, Y., Jang, J., Longpre, S., Lee, H., Yun, S., Shin, S., Kim, S., Thorne, J., & Seo, M. (2024). Prometheus: Inducing fine-grained evaluation capability in language models. In *The Twelfth International Conference on Learning Representations*. <https://arxiv.org/abs/2310.08491>
- Krippendorff, K. (2019). *Content analysis: An introduction to its methodology* (4th ed.). SAGE Publications.
- Landis, J. R., & Koch, G. G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33(1), 159–174. <https://doi.org/10.2307/2529310>
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems* (Vol. 33, pp. 9459–9474). Curran Associates. <https://arxiv.org/abs/2005.11401>
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., & Zhu, C. (2023). G-Eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (pp. 2511–2522). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2023.emnlp-main.153>
- McClure, J., Smyslova, D., Hall, A., & Jiang, S. (2024). Deductive coding’s role in AI vs. human performance. In *Proceedings of the 17th International Conference on Educational Data Mining*. International Educational Data Mining Society. <https://educationaldatamining.org/edm2024/proceedings/2024.EDM-posters.91/index.html>
- O’Connor, C., & Joffe, H. (2020). Inter-coder reliability in qualitative research: Debates and practical guidelines. *International Journal of Qualitative Methods*, 19, Article 1609406919899220. <https://doi.org/10.1177/1609406919899220>
- O’Leary, D. E. (2025). An anchoring effect in large language models. *IEEE Intelligent Systems*, 40(2), 23–26. <https://doi.org/10.1109/MIS.2025.3544939>
- Pangakis, N., Wolken, S., & Fasching, N. (2023). *Automated annotation with generative AI requires validation* (arXiv:2306.00176). arXiv. <https://arxiv.org/abs/2306.00176>
- Rietz, T., & Maedche, A. (2021). Cody: An AI-based system to semi-automate coding for qualitative research. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems* (pp. 1–14). Association for Computing Machinery. <https://doi.org/10.1145/3411764.3445591>
- Tian, K., Mitchell, E., Zhou, A., Sharma, A., Rafailov, R., Yao, H., Finn, C., & Manning, C. (2023). Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (pp. 5433–5442). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2023.emnlp-main.330>
- Wang, P., Li, L., Chen, L., Cai, Z., Zhu, D., Lin, B., Cao, Y., Kong, L., Liu, Q., Liu, T., & Sui, Z. (2024). Large language models are not fair evaluators. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 9440–9450). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2024.acl-long.511>

- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q. V., & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems* (Vol. 35, pp. 24824–24837). Curran Associates. https://proceedings.neurips.cc/paper_files/paper/2022/hash/9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html
- Williamson, D. M., Xi, X., & Breyer, F. J. (2012). A framework for evaluation and use of automated scoring. *Educational Measurement: Issues and Practice*, 31(1), 2–13. <https://doi.org/10.1111/j.1745-3992.2011.00223.x>
- Xiao, Z., Yuan, X., Liao, Q. V., Abdelghani, R., & Oudeyer, P.-Y. (2023). Supporting qualitative analysis with large language models: Combining codebook with GPT-3 for deductive coding. In *Companion Proceedings of the 28th International Conference on Intelligent User Interfaces* (pp. 75–78). Association for Computing Machinery. <https://doi.org/10.1145/3581754.3584136>
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., & Stoica, I. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems* (Vol. 36). Curran Associates. <https://arxiv.org/abs/2306.05685>